

```
from google.colab import files
```

```
uploaded = files.upload()
```

No file chosen

Upload widget is only available when the cell has been executed in the current browser session. Please rerun

this cell to enable.

Saving ready\_outcome1.csv to ready\_outcome1.csv

```
import pandas as pd
```

```
import io
```

```
outcome1 = pd.read_csv(io.BytesIO(uploaded['ready_outcome1.csv']))
```

```
print(outcome1)
```

	v000	v007	v012	v013	v025	v106	v190	\
0	NG7	2018	40	6	1	4	5	
1	NG7	2018	37	5	1	4	5	
2	NG7	2018	27	3	1	4	5	
3	NG7	2018	27	3	1	3	5	
4	NG7	2018	24	2	1	3	5	
...	...	...	...	...	...	...	...	
120648	ZW7	2015	39	5	0	2	3	
120649	ZW7	2015	32	4	0	3	3	
120650	ZW7	2015	24	2	0	2	2	
120651	ZW7	2015	40	6	0	2	3	
120652	ZW7	2015	40	6	0	2	3	

	v501	v501a	v501b	...	\
0	living with partner	Unmarried	Never married	...	
1	married	Married	Married	...	
2	married	Married	Married	...	
3	married	Married	Married	...	
4	no longer living together/separated	Unmarried	Formerly married	...	
...	...	...	...	...	
120648	married	Married	Married	...	
120649	married	Married	Married	...	
120650	married	Married	Married	...	
120651	married	Married	Married	...	
120652	married	Married	Married	...	

	v000_GH6	v000_KE6	v000_ML7	v000_NG7	v000_ZA7	v000_ZM7	v000_ZW7	\
0	0	0	0	1	0	0	0	
1	0	0	0	1	0	0	0	
2	0	0	0	1	0	0	0	
3	0	0	0	1	0	0	0	
4	0	0	0	1	0	0	0	
...	...	...	...	...	...	...	...	
120648	0	0	0	0	0	0	1	
120649	0	0	0	0	0	0	1	
120650	0	0	0	0	0	0	1	
120651	0	0	0	0	0	0	1	
120652	0	0	0	0	0	0	1	

	v501b_Formerly married	v501b_Married	v501b_Never married
0	0	0	1
1	0	1	0
2	0	1	0
3	0	1	0
4	1	0	0
...	...	...	...
120648	0	1	0
120649	0	1	0
120650	0	1	0
120651	0	1	0
120652	0	1	0

[120653 rows x 33 columns]

```
from sklearn.model_selection import train_test_split
```

```
outcome1_train, outcome1_test = train_test_split(
    outcome1, test_size=0.20, stratify=outcome1[['v000', 'child_loss']])
```

```
x_train = outcome1_train[['v000_A07', 'v000_EG6', 'v000_ET7', 'v000_GA6', 'v000_GH6', 'v000_KE6', 'v000_ML7', 'v000_NG7', 'v000_ZA7',
    'v000_ZM7', 'v000_ZW7', 'v013', 'v025', 'v106', 'v190', 'v501b_Formerly married', 'v501b_Married',
    'v501b_Never married', 'first.sex', 'v228', 'smoke']]
```

```
y_train = outcome1_train['lost_child']
```

```
y2_train = outcome1_train['child_loss']
```

```
x_test = outcome1_test[['v000_A07', 'v000_EG6', 'v000_ET7', 'v000_GA6', 'v000_GH6', 'v000_KE6', 'v000_ML7', 'v000_NG7', 'v000_ZA7',
```

```

'v000_ZM7', 'v000_ZW7', 'v013', 'v025', 'v106', 'v190', 'v501b_Formerly married', 'v501b_Married',
'v501b_Never married', 'first.sex', 'v228', 'smoke']]
y_test = outcome1_test['lost_child']
y2_test = outcome1_test['child_loss']

print(outcome1_train.shape, outcome1_test.shape)

print(x_train.shape, x_test.shape)


(96522, 33) (24131, 33)
(96522, 21) (24131, 21)

# roc curve and auc
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import roc_curve
from sklearn.metrics import roc_auc_score
from matplotlib import pyplot

# fit a model
model = LogisticRegression(solver='lbfgs')
model.fit(x_train, y2_train)
# predict probabilities
lr_probs = model.predict_proba(x_test)
# keep probabilities for the positive outcome only
lr_probs = lr_probs[:, 1]
# calculate scores
lr_auc = roc_auc_score(y2_test, lr_probs)
# summarize scores
print('Logistic: ROC AUC=%.3f' % (lr_auc))
# calculate roc curves
lr_fpr, lr_tpr, _ = roc_curve(y2_test, lr_probs)
# plot the roc curve for the model
pyplot.plot(lr_fpr, lr_tpr, marker='.', label='Logistic')
# axis labels
pyplot.xlabel('False Positive Rate')
pyplot.ylabel('True Positive Rate')
# show the legend
pyplot.legend()
# show the plot
pyplot.show()

```

Logistic: ROC AUC=0.743



```

from sklearn.neighbors import KNeighborsClassifier
knn_train_accuracy = []
knn_test_accuracy = []

knn_train2_accuracy = []
knn_test2_accuracy = []

neighbors_settings = range(1, 20)

for n_neighbors in neighbors_settings:
    clf = KNeighborsClassifier(n_neighbors = n_neighbors)
    clf2 = KNeighborsClassifier(n_neighbors = n_neighbors)
    clf.fit(x_train, y_train)
    clf2.fit(x_train, y2_train)

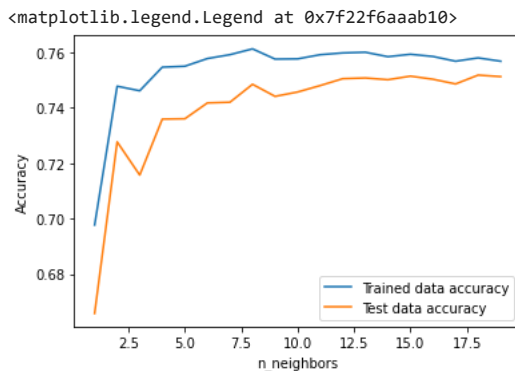
    knn_train_accuracy.append(clf.score(x_train, y_train))
    knn_test_accuracy.append(clf.score(x_test, y_test))

```


```
knn_train2_accuracy.append(clf.score(x_train, y2_train))
knn_test2_accuracy.append(clf.score(x_test, y2_test))
```

```
%matplotlib inline
import matplotlib.pyplot as plt
```

```
plt.plot(neighbors_settings, knn_train_accuracy, label= "Trained data accuracy")
plt.plot(neighbors_settings, knn_test_accuracy, label= "Test data accuracy")
plt.ylabel("Accuracy")
plt.xlabel("n_neighbors")
plt.legend()
```



```
plt.plot(neighbors_settings, knn_train2_accuracy, label= "Trained data accuracy")
plt.plot(neighbors_settings, knn_test2_accuracy, label= "Test data accuracy")
plt.ylabel("Accuracy")
plt.xlabel("n_neighbors")
plt.legend()
```



```
print (knn_train_accuracy)
```

```
print (knn_train2_accuracy)
```

[0.6977580240774124, 0.7478087897059738, 0.7461614968608193, 0.7546776900602972, 0.755092206957999, 0.7577875897681357, 0.7591844346366  
[0.686672468452788, 0.75077184476078, 0.7307764033070181, 0.7444831230185864, 0.7466795134787924, 0.746938521787779, 0.747860591367771,

```
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.naive_bayes import MultinomialNB
from sklearn.pipeline import make_pipeline
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import cross_val_predict
```

```
model = MultinomialNB()
model.fit(x_train, y_train)
labels = model.predict(x_test)
nb_train_accuracy = []
nb_test_accuracy = []
nb_train_accuracy.append(model.score(x_train, y_train))
nb_test_accuracy.append(model.score(x_test, y_test))
print(nb_train_accuracy, nb_test_accuracy)
```

```

model_2 = MultinomialNB()
model_2.fit(x_train, y2_train)
labels_2 = model_2.predict(x_test)
nb2_train_accuracy = []
nb2_test_accuracy = []
nb2_train_accuracy.append(model_2.score(x_train, y2_train))
nb2_test_accuracy.append(model_2.score(x_test, y2_test))
print(nb2_train_accuracy, nb2_test_accuracy)

cv_train = cross_val_score(model, x_train, y_train, cv = 10, scoring = "accuracy")
cv_test = cross_val_score(model, x_test, y_test, cv = 10, scoring = "accuracy")

cv_train2 = cross_val_score(model_2, x_train, y2_train, cv = 10, scoring = "accuracy")
cv_test2 = cross_val_score(model_2, x_test, y2_test, cv = 10, scoring = "accuracy")

print("Percent accuracy within each fold (First): \n")
print(cv_train, "\n", cv_test)
print("\nMean & SD accuracy: \n")
print(cv_train.mean(), cv_test.mean())
print(cv_train.std(), cv_test.std())

print("\nPercent accuracy within each fold (Second): \n")
print(cv_train2, "\n", cv_test2)
print("\nMean & SD accuracy: \n")
print(cv_train2.mean(), cv_test2.mean())
print(cv_train2.std(), cv_test2.std())

cv_labels = cross_val_predict(model, x_test, y_test, cv=10)
cv_labels2 = cross_val_predict(model_2, x_test, y2_test, cv=10)

accuracy = pd.DataFrame()
accuracy2 = pd.DataFrame()

accuracy = accuracy.append([["Naive Baye's", (cv_train.mean()), (cv_test.mean())]])
accuracy2 = accuracy2.append([["Naive Baye's", (cv_train2.mean()), (cv_test2.mean())]])

[0.7540975114481673] [0.7543823297832664]
[0.7703425125878038] [0.7687621731382869]
Percent accuracy within each fold (First):

[0.75240858 0.75479126 0.75310816 0.75300456 0.75663075 0.75290095
 0.75476585 0.75435143 0.75486946 0.7538334 ]
[0.75144988 0.75093245 0.75631993 0.75300456 0.75549109 0.75714878
 0.75549109 0.75259014 0.75300456 0.7538334 ]

Mean & SD accuracy:

0.7540664401177554 0.7539265872409563
0.0012051280451909068 0.0019899742075337387

Percent accuracy within each fold (Second):

[0.76898374 0.77105563 0.76543722 0.77186075 0.77289681 0.77092831
 0.77134273 0.76864898 0.77020307 0.77175715]
[0.76677713 0.76336511 0.7708247 0.76916701 0.77786987 0.76543722
 0.77414007 0.76253626 0.76460837 0.76875259]

Mean & SD accuracy:

0.7703114376360753 0.7683478335212023
0.0020363900774159243 0.004633716958841678

from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.naive_bayes import MultinomialNB
from sklearn.pipeline import make_pipeline
from sklearn.metrics import roc_curve
from sklearn.metrics import roc_auc_score
from matplotlib import pyplot

# fit a model
model_2 = MultinomialNB()
model_2.fit(x_train, y2_train)
# predict probabilities
nb_probs = model_2.predict_proba(x_test)
# keep probabilities for the positive outcome only
nb_probs = nb_probs[:, 1]
# calculate scores

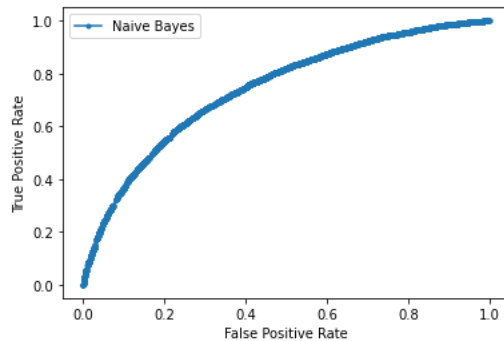
```

```

nb_auc = roc_auc_score(y2_test, nb_probs)
# summarize scores
print('Naive Bayes: ROC AUC=%.3f' % (nb_auc))
# calculate roc curves
nb_fpr, nb_tpr, _ = roc_curve(y2_test, nb_probs)
# plot the roc curve for the model
pyplot.plot(lr_fpr, lr_tpr, marker='.', label='Naive Bayes')
# axis labels
pyplot.xlabel('False Positive Rate')
pyplot.ylabel('True Positive Rate')
# show the legend
pyplot.legend()
# show the plot
pyplot.show()

```

Naive Bayes: ROC AUC=0.729



```

from sklearn.tree import DecisionTreeClassifier

tree = DecisionTreeClassifier(random_state=0)
tree.fit(x_train, y_train)
labels_tree = tree.predict(x_test)

tree_2 = DecisionTreeClassifier(random_state=0)
tree_2.fit(x_train, y2_train)
labels_tree2 = tree_2.predict(x_test)

print("Accuracy on training set: {:.3f}".format(tree.score(x_train, y_train)))
print("Accuracy on test set: {:.3f}".format(tree.score(x_test, y_test)))

print("Accuracy on training set: {:.3f}".format(tree_2.score(x_train, y2_train)))
print("Accuracy on test set: {:.3f}".format(tree_2.score(x_test, y2_test)))

cv_train = cross_val_score(tree, x_train, y_train, cv = 10, scoring = "accuracy")
cv_test = cross_val_score(tree, x_test, y_test, cv = 10, scoring = "accuracy")

cv_train2 = cross_val_score(tree_2, x_train, y2_train, cv = 10, scoring = "accuracy")
cv_test2 = cross_val_score(tree_2, x_test, y2_test, cv = 10, scoring = "accuracy")

print("Percent accuracy within each fold (First): \n")
print(cv_train, "\n", cv_test)
print("\nMean & SD accuracy: \n")
print(cv_train.mean(), cv_test.mean())
print(cv_train.std(), cv_test.std())

print("\nPercent accuracy within each fold (Second): \n")
print(cv_train2, "\n", cv_test2)
print("\nMean & SD accuracy: \n")
print(cv_train2.mean(), cv_test2.mean())
print(cv_train2.std(), cv_test2.std())

cv_labels = cross_val_predict(tree, x_test, y_test, cv=10)
cv_labels2 = cross_val_predict(tree_2, x_test, y2_test, cv=10)

accuracy = accuracy.append(["Decision Tree (Entropy)", (cv_train.mean()), (cv_test.mean())])
accuracy2 = accuracy2.append(["Decision Tree (Entropy)", (cv_train2.mean()), (cv_test2.mean())])

Accuracy on training set: 0.786
Accuracy on test set: 0.754
Accuracy on training set: 0.806
Accuracy on test set: 0.773
Percent accuracy within each fold (First):

```

```
[0.75044028 0.75427328 0.75321177 0.75145048 0.75393701 0.7538334
 0.75393701 0.75082884 0.75238293 0.75435143]
[0.73363712 0.73518442 0.71653543 0.72689598 0.72813925 0.74264401
 0.72440945 0.7273104 0.73062578 0.7273104 ]
```

Mean & SD accuracy:

```
0.7528646424206326 0.7292692234928795
0.0014073972366819359 0.0065965379951465125
```

Percent accuracy within each fold (Second):

```
[0.77250596 0.77240236 0.77154994 0.7708247 0.7745545 0.77393286
 0.77496892 0.77310402 0.76978864 0.77186075]
[0.75476388 0.74595939 0.74015748 0.75341898 0.75093245 0.76004973
 0.74513054 0.74513054 0.75134687 0.74678823]
```


Mean & SD accuracy:

```
0.7725492653510293 0.7493678092052474
0.0015552839454131431 0.005522756117813724
```

```
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.metrics import roc_curve
from sklearn.metrics import roc_auc_score
from matplotlib import pyplot
from sklearn.tree import DecisionTreeClassifier
```

```
# fit a model
model_2 = DecisionTreeClassifier(random_state=0)
model_2.fit(x_train, y2_train)
# predict probabilities
dtc_probs = model_2.predict_proba(x_test)
# keep probabilities for the positive outcome only
dtc_probs = dtc_probs[:, 1]
# calculate scores
dtc_auc = roc_auc_score(y2_test, dtc_probs)
# summarize scores
print('Decision Tree: ROC AUC=%.3f' % (dtc_auc))
# calculate roc curves
dtc_fpr, dtc_tpr, _ = roc_curve(y2_test, dtc_probs)
# plot the roc curve for the model
pyplot.plot(dtc_fpr, dtc_tpr, marker='.', label='Decision Tree')
# axis labels
pyplot.xlabel('False Positive Rate')
pyplot.ylabel('True Positive Rate')
# show the legend
pyplot.legend()
# show the plot
pyplot.show()
```

Decision Tree: ROC AUC=0.758



```
from sklearn.tree import DecisionTreeClassifier
```

```
clf_gini = DecisionTreeClassifier(criterion = "gini", random_state = 100, max_depth = 10, min_samples_leaf=1)
clf_gini.fit(x_train, y_train)
labels_gini = clf_gini.predict(x_test)
```

```
clf_gini2 = DecisionTreeClassifier(criterion = "gini", random_state = 100, max_depth = 10, min_samples_leaf=1)
clf_gini2.fit(x_train, y2_train)
labels_gini2 = clf_gini2.predict(x_test)
```

```
print("Accuracy on training set: {:.3f}".format(clf_gini.score(x_train, y_train)))
```

```

print("Accuracy on test set: {:.3f}".format(clf_gini.score(x_test, y_test)))

print("Accuracy on training set (Child loss): {:.3f}".format(clf_gini2.score(x_train, y2_train)))
print("Accuracy on test set (Child loss): {:.3f}".format(clf_gini2.score(x_test, y2_test)))

cv_train = cross_val_score(clf_gini, x_train, y_train, cv = 10, scoring = "accuracy")
cv_test = cross_val_score(clf_gini, x_test, y_test, cv = 10, scoring = "accuracy")

cv_train2 = cross_val_score(clf_gini2, x_train, y2_train, cv = 10, scoring = "accuracy")
cv_test2 = cross_val_score(clf_gini2, x_test, y2_test, cv = 10, scoring = "accuracy")

print("Percent accuracy within each fold (First): \n")
print(cv_train, "\n", cv_test)
print("\nMean & SD accuracy: \n")
print(cv_train.mean(), cv_test.mean())
print(cv_train.std(), cv_test.std())

print("\nPercent accuracy within each fold (Second): \n")
print(cv_train2, "\n", cv_test2)
print("\nMean & SD accuracy: \n")
print(cv_train2.mean(), cv_test2.mean())
print(cv_train2.std(), cv_test2.std())

cv_labels = cross_val_predict(clf_gini, x_test, y_test, cv=10)
cv_labels2 = cross_val_predict(clf_gini2, x_test, y2_test, cv=10)

accuracy = accuracy.append(["Decision Tree (Gini)", (cv_train.mean()), (cv_test.mean())])
accuracy2 = accuracy2.append(["Decision Tree (Gini)", (cv_train2.mean()), (cv_test2.mean())])

    Accuracy on training set: 0.761
    Accuracy on test set: 0.755
    Accuracy on training set (Child loss): 0.778
    Accuracy on test set (Child loss): 0.770
    Percent accuracy within each fold (First):

    [0.75282296 0.75675956 0.75663075 0.75393701 0.75466225 0.75714878
     0.75476585 0.75393701 0.7555947  0.7576668 ]
    [0.76056338 0.74968918 0.74305843 0.75549109 0.75051803 0.75010361
     0.74222959 0.74720265 0.74554496 0.74554496]

    Mean & SD accuracy:

    0.755392565452619 0.7489945891678292
    0.0015318416924185215 0.005367389578657666

    Percent accuracy within each fold (Second):

    [0.77126282 0.77302393 0.76968504 0.77134273 0.77393286 0.77776627
     0.77092831 0.76875259 0.76989225 0.77331123]
    [0.77050539 0.76750932 0.7443017  0.77952756 0.76668048 0.76460837
     0.75631993 0.76295068 0.76502279 0.75673436]

    Mean & SD accuracy:

    0.7719898022508265 0.7634160586247305
    0.0024944699016082104 0.008954490148631243

```

```

from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.metrics import roc_curve
from sklearn.metrics import roc_auc_score
from matplotlib import pyplot
from sklearn.tree import DecisionTreeClassifier

# fit a model
model_2 = DecisionTreeClassifier(criterion = "gini", random_state = 100, max_depth = 10, min_samples_leaf=1)
model_2.fit(x_train, y2_train)
# predict probabilities
dtc_probs = model_2.predict_proba(x_test)
# keep probabilities for the positive outcome only
dtc_probs = dtc_probs[:, 1]
# calculate scores
dtc_auc = roc_auc_score(y2_test, dtc_probs)
# summarize scores
print('Decision Tree: ROC AUC=%.3f' % (dtc_auc))
# calculate roc curves
dtc_fpr, dtc_tpr, _ = roc_curve(y2_test, dtc_probs)
# plot the roc curve for the model
pyplot.plot(dtc_fpr, dtc_tpr, marker='.', label='Decision Tree')
# axis labels


```

```

pyplot.xlabel('False Positive Rate')
pyplot.ylabel('True Positive Rate')
# show the legend
pyplot.legend()
# show the plot
pyplot.show()

```

Decision Tree: ROC AUC=0.745



```

from sklearn.ensemble import RandomForestClassifier

forest = RandomForestClassifier(n_estimators=100, random_state=0)
forest.fit(x_train, y_train)
labels_rf = forest.predict(x_test)

forest_2 = RandomForestClassifier(n_estimators=100, random_state=0)
forest_2.fit(x_train, y2_train)
labels_rf2 = forest.predict(x_test)

print("Accuracy on training set: {:.3f}".format(forest.score(x_train, y_train)))
print("Accuracy on test set: {:.3f}".format(forest.score(x_test, y_test)))

print("Accuracy on training set (Child loss): {:.3f}".format(forest_2.score(x_train, y2_train)))
print("Accuracy on test set (Child loss): {:.3f}".format(forest_2.score(x_test, y2_test)))

cv_train = cross_val_score(forest, x_train, y_train, cv = 10, scoring = "accuracy")
cv_test = cross_val_score(forest, x_test, y_test, cv = 10, scoring = "accuracy")

cv_train2 = cross_val_score(forest_2, x_train, y2_train, cv = 10, scoring = "accuracy")
cv_test2 = cross_val_score(forest_2, x_test, y2_test, cv = 10, scoring = "accuracy")

print("Percent accuracy within each fold (First): \n")
print(cv_train, "\n", cv_test)
print("\nMean & SD accuracy: \n")
print(cv_train.mean(), cv_test.mean())
print(cv_train.std(), cv_test.std())

print("\nPercent accuracy within each fold (Second): \n")
print(cv_train2, "\n", cv_test2)
print("\nMean & SD accuracy: \n")
print(cv_train2.mean(), cv_test2.mean())
print(cv_train2.std(), cv_test2.std())

cv_labels = cross_val_predict(forest, x_test, y_test, cv=10)
cv_labels2 = cross_val_predict(forest_2, x_test, y2_test, cv=10)

accuracy = accuracy.append(["Random Forest", (cv_train.mean()), (cv_test.mean())])
accuracy2 = accuracy2.append(["Random Forest", (cv_train2.mean()), (cv_test2.mean())])

```

```

Accuracy on training set: 0.786
Accuracy on test set: 0.755
Accuracy on training set (Child loss): 0.806
Accuracy on test set (Child loss): 0.773
Percent accuracy within each fold (First):

```

```

[0.7518906 0.75675956 0.75486946 0.75290095 0.75497306 0.75486946
 0.75663075 0.75300456 0.75642354 0.75414422]
[0.74316487 0.73684211 0.73186904 0.7310402 0.73559884 0.75341898
 0.73435557 0.73352673 0.73642768 0.73352673]

```

```

Mean & SD accuracy:

```

```

0.7546466157256468 0.7369770756373153
0.001592367488739728 0.006347166156062973

```



Percent accuracy within each fold (Second):

```
[0.77074485 0.77426707 0.77041028 0.77020307 0.77673021 0.77755906
 0.77818069 0.77507252 0.77527973 0.77175715]
[0.75724938 0.74761707 0.75051803 0.76212184 0.75217571 0.76419395
 0.74968918 0.75176129 0.75176129 0.74968918]
```


Mean & SD accuracy:

```
0.7740204619569686 0.7536776937679808
0.002887619475658153 0.005317746940072504
```

```
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.metrics import roc_curve
from sklearn.metrics import roc_auc_score
from matplotlib import pyplot
from sklearn.ensemble import RandomForestClassifier

# fit a model
model_2 = RandomForestClassifier(n_estimators=100, random_state=0)
model_2.fit(x_train, y2_train)
# predict probabilities
dttc_probs = model_2.predict_proba(x_test)
# keep probabilities for the positive outcome only
dttc_probs = dttc_probs[:, 1]
# calculate scores
dttc_auc = roc_auc_score(y2_test, dttc_probs)
# summarize scores
print('Random Forest: ROC AUC=%.3f' % (dttc_auc))
# calculate roc curves
dttc_fpr, dttc_tpr, _ = roc_curve(y2_test, dttc_probs)
# plot the roc curve for the model
pyplot.plot(dttc_fpr, dttc_tpr, marker='.', label='Random Forest')
# axis labels
pyplot.xlabel('False Positive Rate')
pyplot.ylabel('True Positive Rate')
# show the legend
pyplot.legend()
# show the plot
pyplot.show()
```

Decision Tree: ROC AUC=0.766



```
from sklearn.ensemble import GradientBoostingClassifier

gbdt = GradientBoostingClassifier(random_state=0, max_depth = 10)
gbdt.fit(x_train, y_train)
labels_gbdt = gbdt.predict(x_test)

gbdt_2 = GradientBoostingClassifier(random_state=0, max_depth = 10)
gbdt_2.fit(x_train, y2_train)
labels_gbdt2 = gbdt_2.predict(x_test)

print("Accuracy on training set: {:.3f}".format(gbdt.score(x_train, y_train)))
print("Accuracy on test set: {:.3f}".format(gbdt.score(x_test, y_test)))

print("Accuracy on training set (Child loss): {:.3f}".format(gbdt_2.score(x_train, y2_train)))
print("Accuracy on test set (Child loss): {:.3f}".format(gbdt_2.score(x_test, y2_test)))

cv_train = cross_val_score(gbdt, x_train, y_train, cv = 10, scoring = "accuracy")
cv_test = cross_val_score(gbdt, x_test, y_test, cv = 10, scoring = "accuracy")
```

```

cv_train2 = cross_val_score(gbrt_2, x_train, y2_train, cv = 10, scoring = accuracy )
cv_test2 = cross_val_score(gbrt_2, x_test, y2_test, cv = 10, scoring = "accuracy")

print("Percent accuracy within each fold (First): \n")
print(cv_train, "\n", cv_test)
print("\nMean & SD accuracy: \n")
print(cv_train.mean(), cv_test.mean())
print(cv_train.std(), cv_test.std())

print("\nPercent accuracy within each fold (Second): \n")
print(cv_train2, "\n", cv_test2)
print("\nMean & SD accuracy: \n")
print(cv_train2.mean(), cv_test2.mean())
print(cv_train2.std(), cv_test2.std())

cv_labels = cross_val_predict(gbrt, x_test, y_test, cv=10)
cv_labels2 = cross_val_predict(gbrt_2, x_test, y2_test, cv=10)

accuracy = accuracy.append(["Gradient Boost", (cv_train.mean()), (cv_test.mean())])
accuracy2 = accuracy2.append(["Gradient Boost", (cv_train2.mean()), (cv_test2.mean())])

```

```

➦ Accuracy on training set: 0.782
Accuracy on test set: 0.758
Accuracy on training set (Child loss): 0.798
Accuracy on test set (Child loss): 0.774
Percent accuracy within each fold (First):

[0.75479126 0.75945302 0.75880646 0.75528388 0.75714878 0.7595317
 0.7595317 0.75549109 0.75787402 0.75797762]
[0.73943662 0.74222959 0.73228346 0.73518442 0.73849979 0.74844592
 0.73062578 0.74098632 0.73974306 0.73477    ]

```

Mean & SD accuracy:

```

0.7575889531260217 0.7382204957886565
0.0017456579478918016 0.0049619438942587885

```

Percent accuracy within each fold (Second):

```

[0.77374909 0.77654615 0.77289681 0.7727932 0.77911314 0.78201409
 0.77704103 0.77445089 0.77486531 0.77807708]
[0.76636288 0.76626606 0.75466225 0.76336511 0.75507667 0.76792375
 0.75507667 0.75839204 0.76170742 0.75466225]

```

Mean & SD accuracy:

```

0.776154679908361 0.7603495083761633
0.0028281639056204184 0.0051342023407811325

```


```

from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.metrics import roc_curve
from sklearn.metrics import roc_auc_score
from matplotlib import pyplot
from sklearn.ensemble import GradientBoostingClassifier

# fit a model
model_2 = GradientBoostingClassifier(random_state=0, max_depth = 10)
model_2.fit(x_train, y2_train)
# predict probabilities
dtc_probs = model_2.predict_proba(x_test)
# keep probabilities for the positive outcome only
dtc_probs = dtc_probs[:, 1]
# calculate scores
dtc_auc = roc_auc_score(y2_test, dtc_probs)
# summarize scores
print('Decision Tree: ROC AUC=%.3f' % (dtc_auc))
# calculate roc curves
dtc_fpr, dtc_tpr, _ = roc_curve(y2_test, dtc_probs)
# plot the roc curve for the model
pyplot.plot(dtc_fpr, dtc_tpr, marker='.', label='Gradient Boosting')
# axis labels
pyplot.xlabel('False Positive Rate')
pyplot.ylabel('True Positive Rate')
# show the legend
pyplot.legend()
# show the plot
pyplot.show()

```

Decision Tree: ROC AUC=0.769



```
from sklearn.svm import SVC
```

```
svc = SVC()
svc.fit(x_train, y_train)
labels_svc = svc.predict(x_test)
```

```
svc_2 = SVC()
svc_2.fit(x_train, y2_train)
labels_svc2 = svc_2.predict(x_test)
```

```
print("Accuracy on training set: {:.3f}".format(svc.score(x_train, y_train)))
print("Accuracy on test set: {:.3f}".format(svc.score(x_test, y_test)))
```

```
print("Accuracy on training set (Child loss): {:.3f}".format(svc_2.score(x_train, y2_train)))
print("Accuracy on test set (Child loss): {:.3f}".format(svc_2.score(x_test, y2_test)))
```

```
cv_train = cross_val_score(svc, x_train, y_train, cv = 10, scoring = "accuracy")
cv_test = cross_val_score(svc, x_test, y_test, cv = 10, scoring = "accuracy")
```

```
cv_train2 = cross_val_score(svc_2, x_train, y2_train, cv = 10, scoring = "accuracy")
cv_test2 = cross_val_score(svc_2, x_test, y2_test, cv = 10, scoring = "accuracy")
```

```
print("Percent accuracy within each fold (First): \n")
print(cv_train, "\n", cv_test)
print("\nMean & SD accuracy: \n")
print(cv_train.mean(), cv_test.mean())
print(cv_train.std(), cv_test.std())
```

```
print("\nPercent accuracy within each fold (Second): \n")
print(cv_train2, "\n", cv_test2)
print("\nMean & SD accuracy: \n")
print(cv_train2.mean(), cv_test2.mean())
print(cv_train2.std(), cv_test2.std())
```

```
cv_labels = cross_val_predict(svc, x_test, y_test, cv=10)
cv_labels2 = cross_val_predict(svc_2, x_test, y2_test, cv=10)
```

```
accuracy = accuracy.append(["SVM", (cv_train.mean()), (cv_test.mean())])
accuracy2 = accuracy2.append(["SVM", (cv_train2.mean()), (cv_test2.mean())])
```

```
Accuracy on training set: 0.755
Accuracy on test set: 0.755
Accuracy on training set (Child loss): 0.769
Accuracy on test set (Child loss): 0.768
```

```
-----
NameError                                Traceback (most recent call last)
<ipython-input-4-f4c30bb4500c> in <module>
    15 print("Accuracy on test set (Child loss): {:.3f}".format(svc_2.score(x_test, y2_test)))
    16
--> 17 cv_train = cross_val_score(svc, x_train, y_train, cv = 10, scoring = "accuracy")
    18 cv_test = cross_val_score(svc, x_test, y_test, cv = 10, scoring = "accuracy")
    19
```

```
NameError: name 'cross_val_score' is not defined
```

SEARCH STACK OVERFLOW

```
cv_train = cross_val_score(svc, x_train, y_train, cv = 10, scoring = "accuracy")
cv_test = cross_val_score(svc, x_test, y_test, cv = 10, scoring = "accuracy")
```

```
cv_train2 = cross_val_score(svc_2, x_train, y2_train, cv = 10, scoring = "accuracy")
cv_test2 = cross_val_score(svc_2, x_test, y2_test, cv = 10, scoring = "accuracy")
```

```

print("Percent accuracy within each fold (First): \n")
print(cv_train, "\n", cv_test)
print("\nMean & SD accuracy: \n")
print(cv_train.mean(), cv_test.mean())
print(cv_train.std(), cv_test.std())

print("\nPercent accuracy within each fold (Second): \n")
print(cv_train2, "\n", cv_test2)
print("\nMean & SD accuracy: \n")
print(cv_train2.mean(), cv_test2.mean())
print(cv_train2.std(), cv_test2.std())

cv_labels = cross_val_predict(svc, x_test, y_test, cv=10)
cv_labels2 = cross_val_predict(svc_2, x_test, y2_test, cv=10)

accuracy = accuracy.append(["SVM", (cv_train.mean()), (cv_test.mean())])
accuracy2 = accuracy2.append(["SVM", (cv_train2.mean()), (cv_test2.mean())])

Percent accuracy within each fold (First):

[0.75427328 0.75530923 0.75435143 0.75528388 0.75528388 0.75590551
 0.75600912 0.75507667 0.75777041 0.75497306]
[0.75144988 0.75217571 0.75217571 0.75259014 0.75176129 0.75341898
 0.75300456 0.75134687 0.75259014 0.75217571]

Mean & SD accuracy:

0.7554236470937297 0.752268899715055
0.0009450310292319527 0.0006224884138372644

Percent accuracy within each fold (Second):

[0.7677406 0.76680825 0.76844177 0.76719851 0.77206797 0.76792375
 0.76906341 0.76875259 0.76947783 0.76989225]
[0.76594863 0.75839204 0.77248239 0.7670949 0.77041028 0.76626606
 0.76999586 0.76336511 0.76295068 0.76585164]

Mean & SD accuracy:

0.768736691371271 0.7662757584486957
0.0014461655059533117 0.003899930540227518

```

```

from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.metrics import roc_curve
from sklearn.metrics import roc_auc_score
from matplotlib import pyplot
from sklearn.svm import SVC

# fit a model
model_2 = SVC()
model_2.fit(x_train, y2_train)
# predict probabilities
dtc_probs = model_2.predict_proba(x_test)
# keep probabilities for the positive outcome only
dtc_probs = dtc_probs[:, 1]
# calculate scores
dtc_auc = roc_auc_score(y2_test, dtc_probs)
# summarize scores
print('Decision Tree: ROC AUC=%.3f' % (dtc_auc))
# calculate roc curves
dtc_fpr, dtc_tpr, _ = roc_curve(y2_test, dtc_probs)
# plot the roc curve for the model
pyplot.plot(dtc_fpr, dtc_tpr, marker='.', label='Decision Tree')
# axis labels
pyplot.xlabel('False Positive Rate')
pyplot.ylabel('True Positive Rate')
# show the legend
pyplot.legend()
# show the plot
pyplot.show()

```

```

-----
AttributeError                                Traceback (most recent call last)
<ipython-input-12-695de4de1df7> in <module>
      9 model_2.fit(x_train, y2_train)
     10 # predict probabilities
--> 11 dtc_probs = model_2.predict_proba(x_test)
     12 # keep probabilities for the positive outcome only
     13 dtc_probs = dtc_probs[:, 1]

-----
1 frames
/usr/local/lib/python3.8/dist-packages/sklearn/svm/_base.py in _check_proba(self)
    798     def _check_proba(self):
    799         if not self.probability:

from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import cross_val_predict

clf = KNeighborsClassifier(n_neighbors = 8)
clf2 = KNeighborsClassifier(n_neighbors = 2)
clf.fit(x_train, y_train)
clf2.fit(x_train, y2_train)

print("Accuracy on training set: {:.3f}".format(clf.score(x_train, y_train)))
print("Accuracy on test set: {:.3f}".format(clf.score(x_test, y_test)))

print("Accuracy on training set (Child loss): {:.3f}".format(clf2.score(x_train, y2_train)))
print("Accuracy on test set (Child loss): {:.3f}".format(clf2.score(x_test, y2_test)))

cv_train = cross_val_score(clf, x_train, y_train, cv = 10, scoring = "accuracy")
cv_test = cross_val_score(clf, x_test, y_test, cv = 10, scoring = "accuracy")

cv_train2 = cross_val_score(clf2, x_train, y2_train, cv = 10, scoring = "accuracy")
cv_test2 = cross_val_score(clf2, x_test, y2_test, cv = 10, scoring = "accuracy")

print("Percent accuracy within each fold (First): \n")
print(cv_train, "\n", cv_test)
print("\nMean & SD accuracy: \n")
print(cv_train.mean(), cv_test.mean())
print(cv_train.std(), cv_test.std())

print("\nPercent accuracy within each fold (Second): \n")
print(cv_train2, "\n", cv_test2)
print("\nMean & SD accuracy: \n")
print(cv_train2.mean(), cv_test2.mean())
print(cv_train2.std(), cv_test2.std())

cv_labels = cross_val_predict(clf, x_test, y_test, cv=10)
cv_labels2 = cross_val_predict(clf2, x_test, y2_test, cv=10)

accuracy = accuracy.append(["KNN", (cv_train.mean()), (cv_test.mean())])
accuracy2 = accuracy2.append(["KNN", (cv_train2.mean()), (cv_test2.mean())])

```

Accuracy on training set: 0.759  
 Accuracy on test set: 0.748  
 Accuracy on training set (Child loss): 0.775  
 Accuracy on test set (Child loss): 0.758  
 Percent accuracy within each fold (First):

```
[0.74484616 0.73572982 0.73767095 0.74181517 0.74316204 0.74823871
 0.74005387 0.74564857 0.74254041 0.74088272]
[0.74606462 0.7405719 0.74140075 0.74264401 0.75217571 0.74388728
 0.74761707 0.74595939 0.7405719 0.74264401]
```

Mean & SD accuracy:

```
0.7420588419396318 0.7443536649555311
0.0035364894759353153 0.003474572792013889
```

Percent accuracy within each fold (Second):

```
[0.7518906 0.74743603 0.7555947 0.7537298 0.75580191 0.75435143
 0.75145048 0.7594281 0.75528388 0.75300456]
[0.74606462 0.75051803 0.7480315 0.74595939 0.75051803 0.75134687
 0.75922089 0.74098632 0.74388728 0.74720265]
```

Mean & SD accuracy:

```
0.7537971474654559 0.748373557205842
0.0030284279033088975 0.004722795807931974
```

```
-----
NameError                                Traceback (most recent call last)
<ipython-input-8-d594213c5a2b> in <module>
    36 cv_labels2 = cross_val_predict(clf2, x_test, y2_test, cv=10)
    37
--> 38 accuracy = accuracy.append(["KNN", (cv_train.mean()), (cv_test.mean())])
    39 accuracy2 = accuracy2.append(["KNN", (cv_train2.mean()), (cv_test2.mean())])

NameError: name 'accuracy' is not defined
```

