

```
from google.colab import files
```

```
uploaded = files.upload()
```

ready.outcome2.csv

- **ready.outcome2.csv**(text/csv) - 9167835 bytes, last modified: 11/19/2022 - 100% done
Saving ready.outcome2.csv to ready.outcome2.csv

```
import pandas as pd
import io
```

```
outcome2 = pd.read_csv(io.BytesIO(uploaded['ready.outcome2.csv']))
print(outcome2)
```

	v000	v007	v012	v013	v025	v106	v190	\
0	NG7	2018	40	6	1	4	5	
1	NG7	2018	37	5	1	4	5	
2	NG7	2018	27	3	1	4	5	
3	NG7	2018	27	3	1	3	5	
4	NG7	2018	24	2	1	3	5	
...	
88392	ZW7	2015	31	4	0	2	2	
88393	ZW7	2015	39	5	0	2	3	
88394	ZW7	2015	32	4	0	3	3	
88395	ZW7	2015	24	2	0	2	2	
88396	ZW7	2015	40	6	0	2	3	

	v501	v501a	v501b	...	\
0	living with partner	Unmarried	Never married	...	
1	married	Married	Married	...	
2	married	Married	Married	...	
3	married	Married	Married	...	
4	no longer living together/separated	Unmarried	Formerly married	...	
...	
88392	married	Married	Married	...	
88393	married	Married	Married	...	
88394	married	Married	Married	...	
88395	married	Married	Married	...	
88396	married	Married	Married	...	

	v000_GH6	v000_KE6	v000_ML7	v000_NG7	v000_ZA7	v000_ZM7	v000_ZW7	\
0	0	0	0	1	0	0	0	
1	0	0	0	1	0	0	0	
2	0	0	0	1	0	0	0	
3	0	0	0	1	0	0	0	
4	0	0	0	1	0	0	0	
...	
88392	0	0	0	0	0	0	1	
88393	0	0	0	0	0	0	1	
88394	0	0	0	0	0	0	1	
88395	0	0	0	0	0	0	1	
88396	0	0	0	0	0	0	1	

	v501b_Formerly married	v501b_Married	v501b_Never married
0	0	0	1
1	0	1	0
2	0	1	0
3	0	1	0
4	1	0	0
...
88392	0	1	0
88393	0	1	0
88394	0	1	0
88395	0	1	0
88396	0	1	0

[88397 rows x 33 columns]

```
from sklearn.model_selection import train_test_split
outcome2_train, outcome2_test = train_test_split(
outcome2, test_size=0.35, stratify=outcome2[['v000', 'v481']])
```

```
x_train = outcome2_train[['v000_A07', 'v000_EG6', 'v000_ET7', 'v000_GA6', 'v000_GH6', 'v000_KE6', 'v000_ML7', 'v000_NG7', 'v000_ZA7',
'v000_ZM7', 'v000_ZW7', 'v013', 'v025', 'v106', 'v190', 'v501b_Formerly married', 'v501b_Married',
'v501b_Never married', 'v213', 'v228', 'm14', 'ANC.facility', 'v312', 'm15', 'v401', 'child_loss']]
y_train = outcome2_train['v481']
```

```
x_test = outcome2_test[['v000_A07', 'v000_EG6', 'v000_ET7', 'v000_GA6', 'v000_GH6', 'v000_KE6', 'v000_ML7', 'v000_NG7', 'v000_ZA7',
'v000_ZM7', 'v000_ZW7', 'v013', 'v025', 'v106', 'v190', 'v501b_Formerly married', 'v501b_Married',
'v501b_Never married', 'v213', 'v228', 'm14', 'ANC.facility', 'v312', 'm15', 'v401', 'child_loss']]
y_test = outcome2_test['v481']
```

```

'v000_ZM7', 'v000_ZW7', 'v013', 'v025', 'v106', 'v190', 'v501b_Formerly married', 'v501b_Married',
'v501b_Never married', 'v213', 'v228', 'm14', 'ANC.facility', 'v312', 'm15', 'v401', 'child_loss']]
y_test = outcome2_test['v481']

print(outcome2_train.shape, outcome2_test.shape)
print(x_train.shape, x_test.shape)

(57458, 33) (30939, 33)
(57458, 26) (30939, 26)

# roc curve and auc
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import roc_curve
from sklearn.metrics import roc_auc_score
from matplotlib import pyplot

# fit a model
model = LogisticRegression(solver='lbfgs')
model.fit(x_train, y_train)
# predict probabilities
lr_probs = model.predict_proba(x_test)
# keep probabilities for the positive outcome only
lr_probs = lr_probs[:, 1]
# calculate scores
lr_auc = roc_auc_score(y_test, lr_probs)
# summarize scores
print('Logistic: ROC AUC=%.3f' % (lr_auc))
# calculate roc curves
lr_fpr, lr_tpr, _ = roc_curve(y_test, lr_probs)
# plot the roc curve for the model
pyplot.plot(lr_fpr, lr_tpr, marker='.', label='Logistic')
# axis labels
pyplot.xlabel('False Positive Rate')
pyplot.ylabel('True Positive Rate')
# show the legend
pyplot.legend()
# show the plot
pyplot.show()

```

⚠ /usr/local/lib/python3.8/dist-packages/sklearn/linear_model/_logistic.py:814: ConvergenceWarning: STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

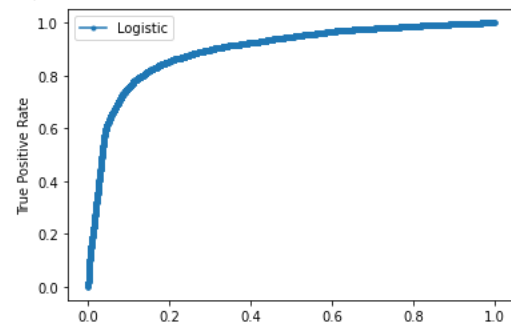
Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```

n_iter_i = _check_optimize_result(
Logistic: ROC AUC=0.897

```



```

from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.naive_bayes import MultinomialNB
from sklearn.pipeline import make_pipeline
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import cross_val_predict

```

```

model = MultinomialNB()
model.fit(x_train, y_train)
labels = model.predict(x_test)
nb_train_accuracy = []
nb_test_accuracy = []
nb_train_accuracy.append(model.score(x_train, y_train))

```

```

nb_test_accuracy.append(model.score(x_test, y_test))
print(nb_train_accuracy, nb_test_accuracy)

cv_train = cross_val_score(model, x_train, y_train, cv = 10, scoring = "accuracy")
cv_test = cross_val_score(model, x_test, y_test, cv = 10, scoring = "accuracy")

print("Percent accuracy within each fold (First): \n")
print(cv_train, "\n", cv_test)
print("\nMean & SD accuracy: \n")
print(cv_train.mean(), cv_test.mean())
print(cv_train.std(), cv_test.std())

cv_labels = cross_val_predict(model, x_test, y_test, cv=10)

accuracy = pd.DataFrame()
accuracy = accuracy.append([[ "Naive Baye's", (cv_train.mean()), (cv_test.mean())]])

[0.9135716523373595] [0.912505252270597]
Percent accuracy within each fold (First):

[0.91106857 0.91385312 0.91176471 0.91420118 0.91646363 0.9096763
 0.91924817 0.90898016 0.91818973 0.91122715]
[0.91758242 0.9114415 0.90723982 0.91305753 0.90950226 0.91338074
 0.91758242 0.91628959 0.91241112 0.90688652]

Mean & SD accuracy:

0.9134672714389879 0.9125373912901725
0.003369901614323646 0.003693114477089309

```

```

from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.naive_bayes import MultinomialNB
from sklearn.pipeline import make_pipeline
from sklearn.metrics import roc_curve
from sklearn.metrics import roc_auc_score
from matplotlib import pyplot

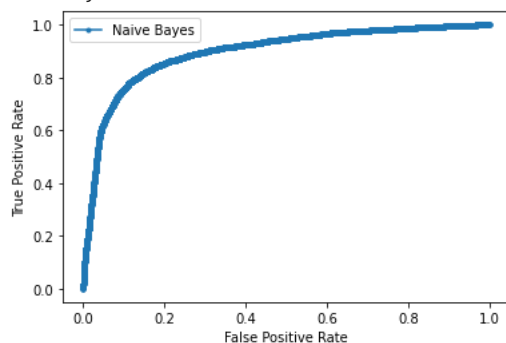
```

```

# fit a model
model_2 = MultinomialNB()
model_2.fit(x_train, y_train)
# predict probabilities
nb_probs = model_2.predict_proba(x_test)
# keep probabilities for the positive outcome only
nb_probs = nb_probs[:, 1]
# calculate scores
nb_auc = roc_auc_score(y_test, nb_probs)
# summarize scores
print('Naive Bayes: ROC AUC=%.3f' % (nb_auc))
# calculate roc curves
nb_fpr, nb_tpr, _ = roc_curve(y_test, nb_probs)
# plot the roc curve for the model
pyplot.plot(nb_fpr, nb_tpr, marker='.', label='Naive Bayes')
# axis labels
pyplot.xlabel('False Positive Rate')
pyplot.ylabel('True Positive Rate')
# show the legend
pyplot.legend()
# show the plot
pyplot.show()

```

Naive Bayes: ROC AUC=0.853



```

from sklearn.tree import DecisionTreeClassifier

tree = DecisionTreeClassifier(random_state=0)
tree.fit(x_train, y_train)
labels_tree = tree.predict(x_test)

print("Accuracy on training set: {:.3f}".format(tree.score(x_train, y_train)))
print("Accuracy on test set: {:.3f}".format(tree.score(x_test, y_test)))

cv_train = cross_val_score(tree, x_train, y_train, cv = 10, scoring = "accuracy")
cv_test = cross_val_score(tree, x_test, y_test, cv = 10, scoring = "accuracy")

print("Percent accuracy within each fold (First): \n")
print(cv_train, "\n", cv_test)
print("\nMean & SD accuracy: \n")
print(cv_train.mean(), cv_test.mean())
print(cv_train.std(), cv_test.std())

cv_labels = cross_val_predict(tree, x_test, y_test, cv=10)

accuracy = accuracy.append(["Decision Tree (Entropy)", (cv_train.mean()), (cv_test.mean())])

```

```

Accuracy on training set: 0.961
Accuracy on test set: 0.898
Percent accuracy within each fold (First):

[0.90184476 0.90288897 0.90219283 0.90236686 0.89784198 0.89906022
 0.90584755 0.8964497  0.89817232 0.89939077]
[0.89010989 0.88946348 0.88914027 0.8804137  0.88623142 0.88849386
 0.89657401 0.88429218 0.89366516 0.88975105]

Mean & SD accuracy:

0.9006055962947851 0.8888135019731985
0.0027302483319281198 0.00429675810289282

```

```

from sklearn.tree import DecisionTreeClassifier

clf_gini = DecisionTreeClassifier(criterion = "gini", random_state = 100, max_depth = 10, min_samples_leaf=1)
clf_gini.fit(x_train, y_train)
labels_gini = clf_gini.predict(x_test)

print("Accuracy on training set: {:.3f}".format(clf_gini.score(x_train, y_train)))
print("Accuracy on test set: {:.3f}".format(clf_gini.score(x_test, y_test)))

cv_train = cross_val_score(clf_gini, x_train, y_train, cv = 10, scoring = "accuracy")
cv_test = cross_val_score(clf_gini, x_test, y_test, cv = 10, scoring = "accuracy")

print("Percent accuracy within each fold (First): \n")
print(cv_train, "\n", cv_test)
print("\nMean & SD accuracy: \n")
print(cv_train.mean(), cv_test.mean())
print(cv_train.std(), cv_test.std())

cv_labels = cross_val_predict(clf_gini, x_test, y_test, cv=10)

accuracy = accuracy.append(["Decision Tree (Gini)", (cv_train.mean()), (cv_test.mean())])

```

```

Accuracy on training set: 0.929
Accuracy on test set: 0.916
Percent accuracy within each fold (First):

[0.91820397 0.91489732 0.91576749 0.92064045 0.91507135 0.9178559
 0.91715976 0.9173338  0.91818973 0.91662315]
[0.91305753 0.91241112 0.91111829 0.90497738 0.91531997 0.91079509
 0.91693601 0.91338074 0.91790562 0.90817976]

Mean & SD accuracy:

0.9171742919053386 0.9124081506063589
0.0016274399443618715 0.0037139120256914947

```

```

from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.metrics import roc_curve
from sklearn.metrics import roc_auc_score
from matplotlib import pyplot
from sklearn.tree import DecisionTreeClassifier


```

```

# fit a model
model_2 = DecisionTreeClassifier(criterion = "gini", random_state = 100, max_depth = 10, min_samples_leaf=1)
model_2.fit(x_train, y_train)
# predict probabilities
dtc_probs = model_2.predict_proba(x_test)
# keep probabilities for the positive outcome only
dtc_probs = dtc_probs[:, 1]
# calculate scores
dtc_auc = roc_auc_score(y_test, dtc_probs)
# summarize scores
print('Decision Tree: ROC AUC=%.3f' % (dtc_auc))
# calculate roc curves
dtc_fpr, dtc_tpr, _ = roc_curve(y_test, dtc_probs)
# plot the roc curve for the model
pyplot.plot(dtc_fpr, dtc_tpr, marker='.', label='Decision Tree')
# axis labels
pyplot.xlabel('False Positive Rate')
pyplot.ylabel('True Positive Rate')
# show the legend
pyplot.legend()
# show the plot
pyplot.show()

```

Decision Tree: ROC AUC=0.886



```

from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.metrics import roc_curve
from sklearn.metrics import roc_auc_score
from matplotlib import pyplot
from sklearn.tree import DecisionTreeClassifier

# fit a model
model_2 = DecisionTreeClassifier(random_state=0)
model_2.fit(x_train, y_train)
# predict probabilities
dtc_probs = model_2.predict_proba(x_test)
# keep probabilities for the positive outcome only
dtc_probs = dtc_probs[:, 1]
# calculate scores
dtc_auc = roc_auc_score(y_test, dtc_probs)
# summarize scores
print('Decision Tree: ROC AUC=%.3f' % (dtc_auc))
# calculate roc curves
dtc_fpr, dtc_tpr, _ = roc_curve(y_test, dtc_probs)
# plot the roc curve for the model
pyplot.plot(dtc_fpr, dtc_tpr, marker='.', label='Decision Tree')
# axis labels
pyplot.xlabel('False Positive Rate')
pyplot.ylabel('True Positive Rate')
# show the legend
pyplot.legend()
# show the plot
pyplot.show()

```

Decision Tree: ROC AUC=0.763



```
from sklearn.ensemble import RandomForestClassifier

forest = RandomForestClassifier(n_estimators=100, random_state=0)
forest.fit(x_train, y_train)
labels_rf = forest.predict(x_test)

print("Accuracy on training set: {:.3f}".format(forest.score(x_train, y_train)))
print("Accuracy on test set: {:.3f}".format(forest.score(x_test, y_test)))

cv_train = cross_val_score(forest, x_train, y_train, cv = 10, scoring = "accuracy")
cv_test = cross_val_score(forest, x_test, y_test, cv = 10, scoring = "accuracy")

print("Percent accuracy within each fold (First): \n")
print(cv_train, "\n", cv_test)
print("\nMean & SD accuracy: \n")
print(cv_train.mean(), cv_test.mean())
print(cv_train.std(), cv_test.std())

cv_labels = cross_val_predict(forest, x_test, y_test, cv=10)

accuracy = accuracy.append(["Random Forest", (cv_train.mean()), (cv_test.mean())])
```

```
Accuracy on training set: 0.961
Accuracy on test set: 0.909
Percent accuracy within each fold (First):

[0.91333101 0.91507135 0.91193874 0.9168117 0.91211277 0.91089454
 0.91350505 0.90880613 0.91314186 0.90966057]
[0.90562379 0.9062702 0.90788623 0.90077569 0.90174531 0.90077569
 0.91241112 0.90820944 0.90982547 0.9033301 ]
```


```
Mean & SD accuracy:
```

```
0.9125273721273389 0.9056853047866913
0.002281606652174675 0.0037844126928809284
```

```
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.metrics import roc_curve
from sklearn.metrics import roc_auc_score
from matplotlib import pyplot
from sklearn.ensemble import RandomForestClassifier

# fit a model
model_2 = RandomForestClassifier(n_estimators=100, random_state=0)
model_2.fit(x_train, y_train)
# predict probabilities
dtc_probs = model_2.predict_proba(x_test)
# keep probabilities for the positive outcome only
dtc_probs = dtc_probs[:, 1]
# calculate scores
dtc_auc = roc_auc_score(y_test, dtc_probs)
# summarize scores
print('Random Forest: ROC AUC=%.3f' % (dtc_auc))
# calculate roc curves
dtc_fpr, dtc_tpr, _ = roc_curve(y_test, dtc_probs)
# plot the roc curve for the model
pyplot.plot(dtc_fpr, dtc_tpr, marker='.', label='Random Forest')
# axis labels
pyplot.xlabel('False Positive Rate')
pyplot.ylabel('True Positive Rate')
# show the legend
pyplot.legend()
# show the plot
pyplot.show()
```

Random Forest: ROC AUC=0.880



```
from sklearn.ensemble import GradientBoostingClassifier

gbrt = GradientBoostingClassifier(random_state=0, max_depth = 10)
gbrt.fit(x_train, y_train)
labels_gbdt = gbdt.predict(x_test)

print("Accuracy on training set: {:.3f}".format(gbdt.score(x_train, y_train)))
print("Accuracy on test set: {:.3f}".format(gbdt.score(x_test, y_test)))

cv_train = cross_val_score(gbdt, x_train, y_train, cv = 10, scoring = "accuracy")
cv_test = cross_val_score(gbdt, x_test, y_test, cv = 10, scoring = "accuracy")

print("Percent accuracy within each fold (First): \n")
print(cv_train, "\n", cv_test)
print("\nMean & SD accuracy: \n")
print(cv_train.mean(), cv_test.mean())
print(cv_train.std(), cv_test.std())

cv_labels = cross_val_predict(gbdt, x_test, y_test, cv=10)

accuracy = accuracy.append(["Gradient Boost", (cv_train.mean()), (cv_test.mean())])
```

Accuracy on training set: 0.952
 Accuracy on test set: 0.912
 Percent accuracy within each fold (First):

```
[0.91646363 0.91350505 0.91367908 0.918378 0.91385312 0.91541942
 0.91872607 0.91141664 0.91679721 0.91470844]
[0.90982547 0.90400776 0.90950226 0.90077569 0.90271493 0.90982547
 0.91758242 0.90885585 0.90982547 0.909473 ]
```


Mean & SD accuracy:

```
0.91529466595296 0.908238832353056
0.002195332906168908 0.0044996540716247405
```

```
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.metrics import roc_curve
from sklearn.metrics import roc_auc_score
from matplotlib import pyplot
from sklearn.ensemble import GradientBoostingClassifier

# fit a model
model_2 = GradientBoostingClassifier(random_state=0, max_depth = 10)
model_2.fit(x_train, y_train)
# predict probabilities
dtt_probs = model_2.predict_proba(x_test)
# keep probabilities for the positive outcome only
dtt_probs = dtt_probs[:, 1]
# calculate scores
dtt_auc = roc_auc_score(y_test, dtt_probs)
# summarize scores
print('Decision Tree: ROC AUC=%.3f' % (dtt_auc))
# calculate roc curves
dtt_fpr, dtt_tpr, _ = roc_curve(y_test, dtt_probs)
# plot the roc curve for the model
pyplot.plot(dtt_fpr, dtt_tpr, marker='.', label='Gradient Boosting')
# axis labels
pyplot.xlabel('False Positive Rate')
pyplot.ylabel('True Positive Rate')
# show the legend
pyplot.legend()
# show the plot
pyplot.show()
```

Decision Tree: ROC AUC=0.899



```
from sklearn.svm import SVC
```

```
svc = SVC()
svc.fit(x_train, y_train)
labels_svc = svc.predict(x_test)
```

```
print("Accuracy on training set: {:.3f}".format(svc.score(x_train, y_train)))
print("Accuracy on test set: {:.3f}".format(svc.score(x_test, y_test)))
```

```
cv_train = cross_val_score(svc, x_train, y_train, cv = 10, scoring = "accuracy")
cv_test = cross_val_score(svc, x_test, y_test, cv = 10, scoring = "accuracy")
```

```
print("Percent accuracy within each fold (First): \n")
print(cv_train, "\n", cv_test)
print("\nMean & SD accuracy: \n")
print(cv_train.mean(), cv_test.mean())
print(cv_train.std(), cv_test.std())
```

```
cv_labels = cross_val_predict(svc, x_test, y_test, cv=10)
```

```
accuracy = accuracy.append(["SVM", (cv_train.mean()), (cv_test.mean())])
```

```
Accuracy on training set: 0.918
Accuracy on test set: 0.917
Percent accuracy within each fold (First):
```

```
[0.91420118 0.91768187 0.9168117 0.91994431 0.92220675 0.91507135
0.9244692 0.91176471 0.91993037 0.9129678 ]
[0.92049127 0.91822883 0.90982547 0.91758242 0.91467356 0.91596639
0.92146089 0.91725921 0.91467356 0.9123828 ]
```

```
Mean & SD accuracy:
```

```
0.9175049233931836 0.9162544402973454
0.0039156881758704775 0.003363507145887208
```

```
from sklearn.neighbors import KNeighborsClassifier
```

```
knn_train_accuracy = []
knn_test_accuracy = []
```

```
neighbors_settings = range(1, 20)
```

```
for n_neighbors in neighbors_settings:
    clf = KNeighborsClassifier(n_neighbors = n_neighbors)
    clf.fit(x_train, y_train)
```

```
knn_train_accuracy.append(clf.score(x_train, y_train))
knn_test_accuracy.append(clf.score(x_test, y_test))
```

```
print (knn_train_accuracy)
print (knn_test_accuracy)
```

```
[0.9477531414250409, 0.9312889414876954, 0.9328727070207804, 0.9285913188763967, 0.928660934943785, 0.9256848480629329, 0.92639841275366,
0.8841591518795048, 0.90222696273312, 0.9045218009631856, 0.9093377290798021, 0.9120527489576263, 0.9114063156533825, 0.913087042244416,
0.9142011801180118, 0.9176818701768187, 0.91681170168117, 0.9199443101994431, 0.9222067502220675, 0.9150713501507135, 0.92446920244692,
0.9117647101176471, 0.9199303701993037, 0.91296780129678, 0.9204912702049127, 0.9182288301822883, 0.9098254700982547, 0.9175824201758242, 0.9146735601467356, 0.9159663901596639,
0.9214608902146089, 0.9172592101725921, 0.9146735601467356, 0.91238280123828]
```


```
%matplotlib inline
```

```
import matplotlib.pyplot as plt
```

```
plt.plot(neighbors_settings, knn_train_accuracy, label= "Trained data accuracy")
```



```
plt.plot(neighbors_settings, knn_test_accuracy, label= "Test data accuracy")
plt.ylabel("Accuracy")
plt.xlabel("n_neighbors")
plt.legend()
```



```
clf = KNeighborsClassifier(n_neighbors = 20)

clf.fit(x_train, y_train)

print("Accuracy on training set: {:.3f}".format(clf.score(x_train, y_train)))
print("Accuracy on test set: {:.3f}".format(clf.score(x_test, y_test)))

cv_train = cross_val_score(clf, x_train, y_train, cv = 10, scoring = "accuracy")
cv_test = cross_val_score(clf, x_test, y_test, cv = 10, scoring = "accuracy")

print("Percent accuracy within each fold (First): \n")
print(cv_train, "\n", cv_test)
print("\nMean & SD accuracy: \n")
print(cv_train.mean(), cv_test.mean())
print(cv_train.std(), cv_test.std())

cv_labels = cross_val_predict(clf, x_test, y_test, cv=10)

accuracy = accuracy.append(["KNN", (cv_train.mean()), (cv_test.mean())])
```

```
Accuracy on training set: 0.920
Accuracy on test set: 0.915
Percent accuracy within each fold (First):

[0.91176471 0.91750783 0.91576749 0.91176471 0.91663766 0.90985033
 0.918378    0.90898016 0.91331593 0.91523064]
[0.90465417 0.90950226 0.90853264 0.91014867 0.90723982 0.91047188
 0.91370394 0.91370394 0.90659341 0.90850307]

Mean & SD accuracy:

0.9139197449801989 0.9093053814825938
0.003096493187079155 0.0027441077092316643
```

✓ 20s completed at 21:36

